

Breast Cancer Data Analysis

```
In [2]: #importing the dataset
df = pd.read_csv('C:/Users/zizhe/Desktop/Leo Zhao/Breast Cancer Prediction/data.csv')
df.head()
```

| nean | concavity_mean | concave points_mean | ... | texture_worst | perimeter_worst | area_worst | smoothness_worst | compactness_worst | concavity_worst | concave points_worst |
|------|----------------|------------------------|-----|---------------|-----------------|------------|------------------|-------------------|-----------------|-------------------------|
| 7760 | 0.3001 | 0.14710 | ... | 17.33 | 184.60 | 2019.0 | 0.1622 | 0.6656 | 0.7119 | 0.2654 |
| 7864 | 0.0869 | 0.07017 | ... | 23.41 | 158.80 | 1956.0 | 0.1238 | 0.1866 | 0.2416 | 0.1860 |
| 5990 | 0.1974 | 0.12790 | ... | 25.53 | 152.50 | 1709.0 | 0.1444 | 0.4245 | 0.4504 | 0.2430 |
| 8390 | 0.2414 | 0.10520 | ... | 26.50 | 98.87 | 567.7 | 0.2098 | 0.8663 | 0.6869 | 0.2575 |
| 3280 | 0.1980 | 0.10430 | ... | 16.67 | 152.20 | 1575.0 | 0.1374 | 0.2050 | 0.4000 | 0.1625 |

Data Preprocessing Part 1

```
In [3]: # dropping unnecessary columns
df.drop([32,'id'],axis=1,inplace=True)
```

```
In [4]: df.head()
```

Out[4]:

| | diagnosis | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean |
|---|-----------|-------------|--------------|----------------|-----------|-----------------|------------------|----------------|---------------------|
| 0 | M | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 |
| 1 | M | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 |
| 2 | M | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 |
| 3 | M | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 |
| 4 | M | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 |

5 rows × 31 columns



```
In [5]: #checking for the missing values  
df.isnull().sum()
```

```
Out[5]: diagnosis          0
radius_mean               0
texture_mean              0
perimeter_mean            0
area_mean                 0
smoothness_mean           0
compactness_mean          0
concavity_mean            0
concave points_mean       0
symmetry_mean             0
fractal_dimension_mean    0
radius_se                 0
texture_se                0
perimeter_se              0
area_se                   0
smoothness_se             0
compactness_se            0
concavity_se              0
concave points_se         0
symmetry_se               0
fractal_dimension_se      0
radius_worst              0
texture_worst             0
perimeter_worst           0
area_worst                0
smoothness_worst          0
compactness_worst         0
concavity_worst           0
concave points_worst      0
symmetry_worst            0
fractal_dimension_worst   0
dtype: int64
```

```
In [6]: #checking the data types of the columns
df.dtypes
```

```
Out[6]: diagnosis      object
radius_mean      float64
texture_mean      float64
perimeter_mean    float64
area_mean         float64
smoothness_mean   float64
compactness_mean  float64
concavity_mean    float64
concave points_mean float64
symmetry_mean     float64
fractal_dimension_mean float64
radius_se         float64
texture_se        float64
perimeter_se      float64
area_se           float64
smoothness_se     float64
compactness_se    float64
concavity_se      float64
concave points_se float64
symmetry_se       float64
fractal_dimension_se float64
radius_worst      float64
texture_worst     float64
perimeter_worst   float64
area_worst        float64
smoothness_worst  float64
compactness_worst float64
concavity_worst   float64
concave points_worst float64
symmetry_worst    float64
fractal_dimension_worst float64
dtype: object
```

```
In [19]: # checking the data description
df.describe()
```

Out[19]:

| | radius_mean | texture_mean | perimeter_mean | area_mean | smoothness_mean | compactness_mean | concavity_mean | concave points_mean | symmr |
|--------------|-------------|--------------|----------------|-------------|-----------------|------------------|----------------|---------------------|-------|
| count | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | |
| mean | 14.127292 | 19.289649 | 91.969033 | 654.889104 | 0.096360 | 0.104341 | 0.088799 | 0.048919 | |
| std | 3.524049 | 4.301036 | 24.298981 | 351.914129 | 0.014064 | 0.052813 | 0.079720 | 0.038803 | |
| min | 6.981000 | 9.710000 | 43.790000 | 143.500000 | 0.052630 | 0.019380 | 0.000000 | 0.000000 | |
| 25% | 11.700000 | 16.170000 | 75.170000 | 420.300000 | 0.086370 | 0.064920 | 0.029560 | 0.020310 | |
| 50% | 13.370000 | 18.840000 | 86.240000 | 551.100000 | 0.095870 | 0.092630 | 0.061540 | 0.033500 | |
| 75% | 15.780000 | 21.800000 | 104.100000 | 782.700000 | 0.105300 | 0.130400 | 0.130700 | 0.074000 | |
| max | 28.110000 | 39.280000 | 188.500000 | 2501.000000 | 0.163400 | 0.345400 | 0.426800 | 0.201200 | |

8 rows × 30 columns

```
In [11]: # Remove any leading or trailing spaces from column names
df.columns = df.columns.strip()
```


```
In [12]: # Check the corrected column names
print(df.columns)
```

```
Index(['diagnosis', 'radius_mean', 'texture_mean', 'perimeter_mean',
       'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
       'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
       'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
       'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
       'fractal_dimension_se', 'radius_worst', 'texture_worst',
       'perimeter_worst', 'area_worst', 'smoothness_worst',
       'compactness_worst', 'concavity_worst', 'concave points_worst',
       'symmetry_worst', 'fractal_dimension_worst'],
      dtype='object')
```

Exploratory Data Analysis


```
In [17]: # bar plot for the number of diagnosis
plt.figure(figsize=(5,5)) #plt.figure(figsize=(5, 5)):
sns.barplot(x=df['diagnosis'].value_counts(),y=df['diagnosis'].value_counts().values)
```


Out[17]: <AxesSubplot:>



```
In [21]: # create a heatmap to check the correlation.  
plt.figure(figsize=(20,20))  
sns.heatmap(df.cor(),annot=True)
```

Out[21]: <AxesSubplot:>





Train Test Split

```
In [22]: #split a DataFrame 'df' into training and testing datasets for machine Learning purposes.

from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(df.drop(['diagnosis']),df['diagnosis'],test_size,random_state=42)
```

Using Decision Tree Classifier

```
In [23]: from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier()
dtree.fit(X_train,y_train)
```

```
Out[23]: DecisionTreeClassifier()
```

```
In [24]: #predicting the diagnosis
y_pred = dtree.predict(X_test)
```

Model Evaluation

```
In [25]: # printing samples from predicted and actual values
print('Predicted values: ',y_pred[:10])
print('Actual values:y_test[:10])
```



```
Predicted values: ['B' 'M' 'M' 'B' 'B' 'M' 'M' 'M' 'B' 'B']
Actual values: 204    B
70      M
131     M
431     B
540     B
567     M
369     M
29      M
81      B
477     B
Name: diagnosis, dtype: object
```

```
In [26]: # model evaluation
print(dtree.score(X_test,y_test))
```

```
0.9415204678362573
```

```
In [ ]: # 0.94 means that the model is making correct predictions for approximately 94% of the samples in the test data.
```

Using logistic regression

```
In [28]: from sklearn.linear_model import LogisticRegression
logmodel = LogisticRegression()
logmodel(X_train,y_train)
```

```
C:\Users\zizhe\New folder\lib\site-packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed to
converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(
```

```
Out[28]: LogisticRegression()
```

```
In [30]: yhat = logmodel.predict(X_test)
```

```
#yhat: This variable stores the predicted labels for the test data. After running this line of code, yhat will be an ar
```

Model Evaluation

```
In [31]: # printing samples from predicted and actual values
print('Predicted values: ',yhat[:10])
print('Actual values: ',y_test[:10])
# If the predicted values closely match the actual values, it's a positive sign that your model is making accurate pred
```

```
Predicted values: ['B' 'M' 'M' 'B' 'B' 'M' 'M' 'M' 'B' 'B']
```

```
Actual values: 204    B
```

```
70      M
```

```
131     M
```

```
431     B
```

```
540     B
```

```
567     M
```

```
369     M
```

```
29      M
```

```
81      B
```

```
477     B
```

```
Name: diagnosis, dtype: object
```

```
In [ ]: #The numbers "204" and "477" you see in the output are actually indices or row numbers from your test dataset. In the c
```

```
In [32]: # model evaluation
print(logmodel.score(X_test,y_test))

#Logmodel.score(X_test, y_test), calculates and prints the accuracy of the logistic regression model (Logmodel) on the
0.9707602339181286
```

```
In [ ]: # Conclusion
```

```
In [ ]: # From both the models we can see that the accuracy is 94% and 97% respectively. But we can see that the recall value f
```